

Robin Willenbrock

Static Detection of Data Races in Interrupt-Driven Software Using Reduced Inter-Procedural Control Flow Graphs

June 28, 2024

supervised by:
Ulrike Engeln

Hamburg University of Technology (TUHH)
Technische Universität Hamburg
Institute for Software Systems
21073 Hamburg

Contents

1	Introduction	1
2	Background	3
2.1	Interrupt-Driven Systems	3
2.2	Shared Resources	4
2.3	Reduced Inter-Procedural Control Flow Graphs (RICFG)	4
2.4	Data Races	6
2.4.1	Detection Techniques	7
2.4.2	Strategies for Preventing Data Races	8
2.5	Static Detection of Data Races in Interrupt-Driven Systems	9
3	Implementation	11
4	Evaluation	21
5	Conclusion	23
	Bibliography	25
	Attachments	29

List of Figures

2.1	Flow-Chart des interruptgesteuerten Systems	3
2.2	Example of Inter-Procedural Control Flow Graph	5
2.3	Example of Reduced Inter-Procedural Control Flow Graph	6
2.4	Simple Example of a Data Race	7
2.5	Example of a Data Race with Enable/Disable ISR Calls	9
2.6	Static Race Detection Approach by [1]	10

1 Introduction

2 Background

2.1 Interrupt-Driven Systems

An interrupt-driven system is an architecture where the flow of execution is changed by unpredictable events in the system, also known as interrupts. Interrupts can be caused by hardware devices, software conditions, or external signals forcing the processor to suspend the current task to execute an interrupt handler or interrupt service routine (ISR). Interrupt-driven systems are used in real-time operating systems, embedded systems, and generally in systems where timely responses are necessary [1].

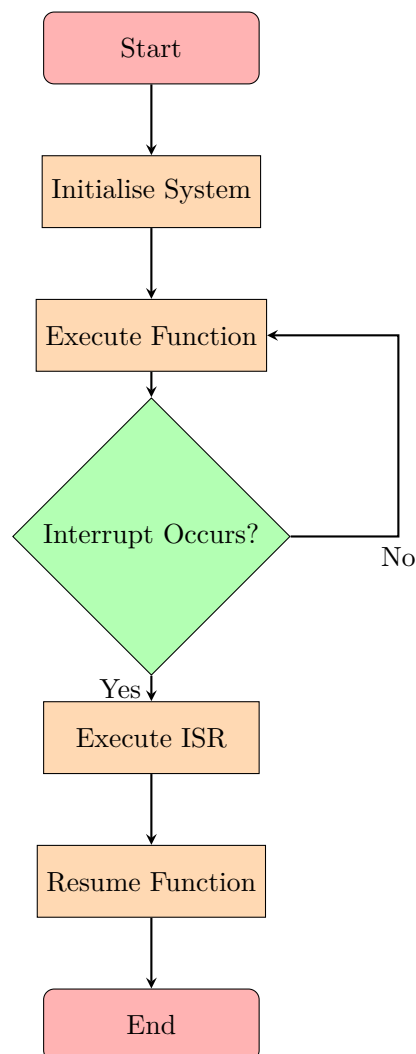


Figure 2.1: Flow-Chart des interruptgesteuerten Systems

In Figure 2.1, a basic execution flow of a simple interrupt-driven system is displayed. The system executes a function as long as no interrupt occurs. When an interrupt occurs, it switches to the ISR, executes it, and then resumes the function executed before the interrupt happened.

The management of the interrupts to maintain the fast responsiveness of the system is the most challenging part of an interrupt-driven system. Interrupts occur in unpredictable ways, so you have to consider every possible execution flow. To ensure the execution of critical interrupts, interrupts are often prioritised, so higher priority events can interrupt lower ones and be handled immediately. When handling an interrupt, the current state of the process is saved, and the context is switched to the ISR [1].

The unpredictability and asynchronous nature of the interrupts present a lot of challenges in designing and implementing an interrupt-driven system. One of the biggest challenges is the correct handling of high-priority interrupts without delaying them substantially. Which needs a sophisticated scheduling and prioritisation mechanism. The execution of the main programme and ISR needs to be handled properly to ensure data integrity. Furthermore, handling context switches, preserving system state, and avoiding deadlocks all contribute to the development of an interrupt-driven system.

2.2 Shared Resources

Shared resources, often referred to as shared memory or shared variables, are data that can be accessed simultaneously by multiple threads or processes. Proper management of these resources is crucial because improper handling can lead to issues like data races, deadlocks, and other synchronisation problems. In interrupt-driven systems, shared resources often involve variables or data structures that are accessed by both the main programme and ISRs. Proper management of shared resources is critical to ensuring data consistency and avoiding conflicts [8]. Proper management of shared resources involves the use of synchronisation mechanisms to coordinate access and ensure data consistency. Mutexes, semaphores, and condition variables are common tools used to control access to shared resources. Mutexes provide mutual exclusion, ensuring that only one thread can access the resource at a time. Semaphores can limit the number of threads accessing the resource simultaneously. Condition variables allow threads to wait for certain conditions to be met before proceeding, facilitating complex synchronisation scenarios [8]. In interrupt driven software, the synchronisation of the shared resources often implies disabling-enabling interrupts [3]. Analysing the management of the shared resources is a large part of the data race analysis, which is further explained later.

2.3 Reduced Inter-Procedural Control Flow Graphs (RICFG)

Control Flow Graphs (CFG) are representations of all possible paths through a programme or a function during its execution. An Inter-Procedural Control Flow Graph (ICFG) adds possible edges between multiple programmes or functions to also show possible control flows between those. A Reduced Inter-Procedural Control Flow Graph

(RICFG) is an optimised version of the ICFG that simplifies the graph to only the necessary information needed for the analysis [2].

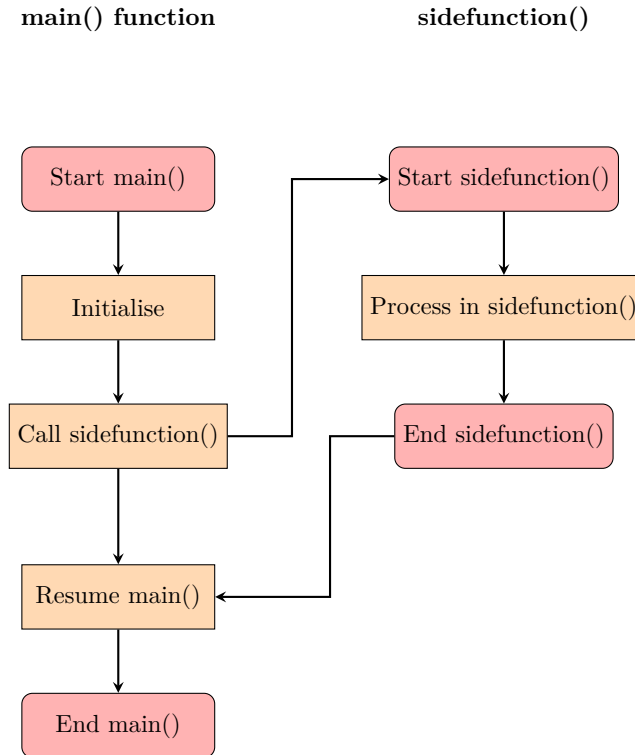


Figure 2.2: Example of Inter-Procedural Control Flow Graph

In Figure 2.2, a simple ICFG is shown. There are two separate linear control flow graphs where the main function calls the sidefunction in its execution. To interpret the flow of the programme correctly, you need to consider the execution of sidefunction() and where it's called. The ICFG combines the two separate CFGs to ensure correct analysis.

There are multiple techniques to reduce the graph, such as node merging, edge contraction, and the elimination of non-important nodes, without losing any information required for the analysis and reducing the complexity of the RICFG. The reduction of the ICFG makes the analysis of large and complex software a lot more efficient. By minimising the amount of data while retaining enough detail, RICFGs are great for static analysis of data races [1].

Node merging is combining nodes that represent redundant control flow paths to reduce the number of nodes in the graph. Edge contraction is simplifying the graph by reducing the number of edges between nodes. It collapses edges, that do not significantly affect the control flow of the graph. [6] The elimination of nodes is the main tool used in this work to reduce the CFG. Eliminating nodes that do not carry any essential information for the applied data analysis significantly reduces the amount of data the

algorithm has to analyze. Overall, these techniques enhance the scalability of static analysis and make it more practical to analyse more complex data [1].

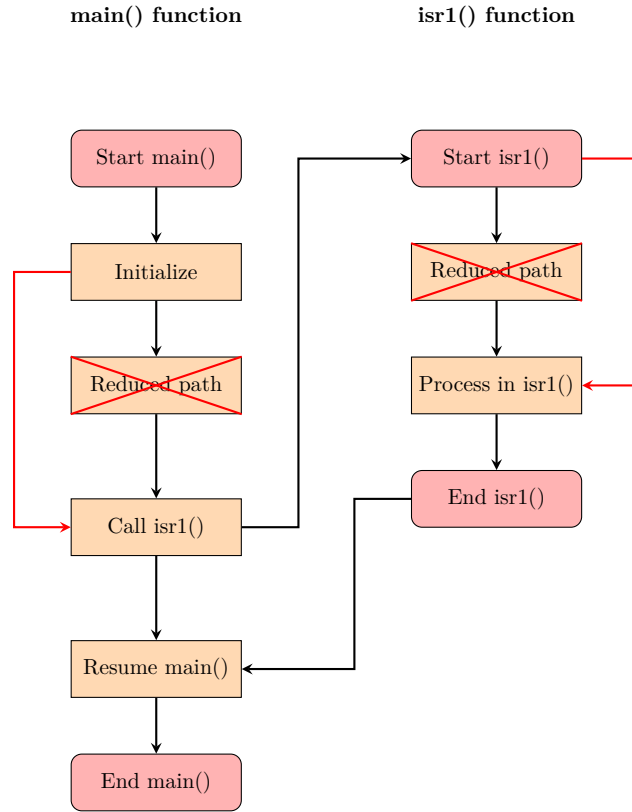


Figure 2.3: Example of Reduced Inter-Procedural Control Flow Graph

Figure 2.3 shows an example of a simple reduction by eliminating nodes that do not carry any important information for the analysis the RICFG is used for.

2.4 Data Races

A data race occurs when two or more functions or threads access a shared resource concurrently, without being ordered by a happens-before relationship, and one of those accesses is a write operation [5]. This can lead to unpredictable behaviour and errors in the system, which makes the detection of data races a critical aspect of concurrent programmes. Without proper synchronisation, a system with multiple threads or functions that use shared data will lead to data races. The outcome of a programme with data races is non-deterministic [5]. The order of execution of operations can vary, which may result in the generation of bugs that are not reproducible or difficult to reproduce.

Algorithm 1: Data Race Example

```

Data: long shared1;
1 Function main ():
2   Variables:
3   unsigned char tmp;
4   Code:
5   tmp ← shared1;
6 Function isr1 ():
7   Code:
8   idlerun();
9   shared1 ← 1;
10  idlerun();

```

Figure 2.4: Simple Example of a Data Race

In Figure 2.4, an example of a simple data race is shown. A global variable `shared1` is initiated and accessed in two different functions `main()` and `isr1()`. Since there are no synchronisation tools used and the operation in `isr1` is a write, there is a data race between line 5 and line 9.

2.4.1 Detection Techniques

Data race detection can be approached by two different analytical methods. Each of those methods provides benefits and challenges.

Static Data Race Detection [1]

Advantages:

- **Comprehensiveness:** Static analysis inspects the code without executing the programme by analysing every possible execution path and interactions that could lead to data races.
- **Early Detection:** Since static analysis does not require execution, it can analyse the code in the development phase, allowing the developer to find issues without deployment.

Disadvantages:

- **False Positives/Negatives:** Static analysis reports all data races that fall under certain conditions. Some of these data races could be very unlikely or even impossible at runtime. On the other hand, due to the approximations and assumptions necessary for tractability, it may miss some races.

- **Complexity in Handling Dynamic Behaviour:** Dynamic behaviours such as pointers or recursion can be challenging to analyse for static approaches, leading to incomplete or inaccurate results.

Dynamic Data Race Detection [4]

Advantages:

- **Precision:** Dynamic analysis tools monitor the actual execution of a programme, identifying data races in real-time. Which results in reducing the number of false positives.
- **Context-Sensitive Detection:** By analysing the actual runtime behaviour, dynamic analysis can understand the context of operations, leading to more accurate detection.

Disadvantages:

- **Performance Overhead:** The analysis in runtime can slow down the application significantly.
- **Coverage:** The effectiveness is heavily dependent on the execution path triggered during the tests. If certain parts of the programme are not passed through in the execution run, they are not analysed.

Both static and dynamic analysis are crucial for a complete analysis of a code. They complement each other's limitations. A combination of both is the best approach to detecting data races most reliable. However, in this work, I am going to focus on the static analysis of data races.

2.4.2 Strategies for Preventing Data Races

Preventing data races requires careful design and implementation of concurrent programs. One effective strategy is to use proper synchronisation mechanisms, such as mutexes, semaphores, and condition variables, to control access to shared data. These mechanisms ensure that only one thread can access the shared data at a time, preventing conflicting operations. Avoiding shared mutable states is another effective strategy, where threads operate on local copies of data instead of shared data, reducing the potential for conflicts. Designing thread-safe data structures and algorithms that inherently manage concurrent access also helps prevent data races, ensuring reliable and predictable programme behaviour [8].

Preventing data races requires careful design and implementation of concurrent programs. Effective strategies for general prevention of data races are synchronisation mechanisms such as mutexes, semaphors, and condition variables, which control access to shared data. Those mechanisms ensure that only one thread can access the shared resource at a time [8]. Since I am focusing on data races in interrupt driven systems,

the main tool to prevent data races is to disable ISRs, which access shared resources in critical areas.

Algorithm 2: Enable/Disable ISR Call Example

```

Data: long shared1
1 Function main():
2   Variables:
3   unsigned char tmp;
4   Code:
5   disable_isr(1);
6   tmp ← shared1;
7   enable_isr(1);
8 Function isr1():
9   Code:
10  idlerun();
11  shared1 ← 1;
12  idlerun();
13 Function isr2():
14  Code:
15  idlerun();
16  int variable1 = 1;
17  idlerun();
  
```

Figure 2.5: Example of a Data Race with Enable/Disable ISR Calls

Figure 2.5 is an example of a disable ISR call that leads to the safe access of the shared data. The main function and *isr1* both access the shared resource *shared1*. Since the read operation in line 6 of the main function is safely accessed by disabling *isr1* in line 5 and enabling it in line 7, a possible data race is prevented.

2.5 Static Detection of Data Races in Interrupt-Driven Systems

The asynchronous nature and concurrent execution of ISRs and the main function introduces significant challenges for data consistency and detecting data races in interrupt-driven systems. Static data race analysis, especially those using RICFGs, are a promising approach to identifying data races without the need for extensive testing and runtime monitoring as in dynamic approaches [1].

The static approach involves the construction of an RICFG for the programme, which includes both the main code and ISRs, and capturing the control flow and potential interaction between them. Analysing the RICFG, shows paths where shared resources

are accessed concurrently without proper synchronisation and indicates potential data races. Integrating the static analysis tool with the development process enables continuous detection of data races during software development, improving the reliability and correctness of interrupt-driven systems [1].

The methodology for static data race detection in interrupt-driven systems involves the following key steps. First, the RICFGs are constructed for the entire programme, including the main code and the ISRs. This involves analysing the control flow and identifying interactions between the main programme and ISRs. Next, the RICFGs are analysed to find potential data races, focusing on paths where concurrent access of shared data is done without proper synchronization. Finally, the developer can use the analysis results to address identified data races in early development processes [1].

Algorithm 3: Static Race Detection

Input: RICFGs of P

Output: potential racing pairs (PR)

```

1 for each  $\langle G_i; G_j \rangle$  in RICFGs do
2   for each  $sv_i \in G_i$  do
3     for each  $sv_j \in G_j$  do
4       if  $sv_i.V == sv_j.V$  and  $(sv_i.A == W$  or  $sv_j.A == W)$  and
          $G_i.pri < G_j.pri$  and  $INTB.get(sv_i).contains(G_j)$  then
5          $PR = PR \cup \{\langle sv_i, sv_j \rangle\};$ 

```

Figure 2.6: Static Race Detection Approach by [1]

The approach by Wang et al. shows a computation of potential data races using RICFGs. By running a depth-first search on the RICFGs it finds the interrupt status of every instruction. If there is a shared resource in both of the analysed RICFGs, at least one of them is a write operation, and the two functions differ in their priority. While the interrupt in this pair is enabled, the two accesses are a potential data race [1]. In the following, I am going to introduce you to the implementation of my static analysis programme based on the static race detection approach of Wang et al..

3 Implementation

In the following I will provide an indepth explanation of my implemenation. For the generation of the input I used GCC. The command `gcc -fdump-tree-cfg` provides a `cfg`-file with all important informations for the intended data race analysis. I have split the explanation of the implementation into the initialization of the basic block class, the parsing of the input, the actual data race analysis and the filtering of the found data races.

Class BasicBlock

Algorithm 4: BasicBlock Class Definition

```

1 Class BasicBlock:
2   Def __init__(self, function_name, number, shared_resources, successors,
    enable_disable_calls, code):
3     self.function_name ← function_name
4     self.number ← number
5     self.shared_resources ← shared_resources
6     self.successors ← successors
7     self.enable_disable_calls ← enable_disable_calls
8     self.code ← code

```

The class **BasicBlock** is a display of all the information nessessary for the data race analysis found in the input. Those informations include the following attributes:

- **function_name:** The function name to which the basic block belongs.
- **number:** The number of the basic block.
- **shared_resources:** All accesses of shared resources within the basic block. The access type (read/write) and also the line number of such calls are saved.
- **successors:** A list of all the successors of each basic block. Important to build all possible paths through the CFG.
- **enable_disable_calls:** All calls that disable or enable an ISR within this basic block and also with the corrsponding line number of those calls to ensure correct order.
- **code:** The lines of code of the block stored with their line number.

Parsing

Algorithm 5: Parse Basic Blocks

Data: file_path, shared_resource_names
Result: Dictionary of BasicBlocks mapping (function name, block number) to BasicBlock instances

```

1 blocks ← empty dictionary
2 current_function ← None
   Input: Path to source file, list of shared resource names
   Output: Dictionary of basic blocks with details
3 for each line in file at file_path do
4   line ← line.strip()
5   Increment line_number
6   if line starts new function then
7     if current block is open then
8       | Save current block to blocks
9     end
10    Update current_function and reset current block details
11  end
12  else if line starts new basic block then
13    if current block is open then
14      | Save current block to blocks
15    end
16    Initialize new block details
17  end
18  Update shared_resources and enable_disable_calls based on line content
19 end
20 if current block is open at end of file then
21   | Save last block to blocks
22 end
23 foreach line indicating successor blocks do
24   | Parse successor details and update block information in blocks
25 end
26 return blocks

```

The function `parse_basic_blocks` starts with the initialisation of a dictionary `blocks` and a variable `current_function`. It opens the file and reads its lines into a list `lines`.

These variables are initialised to collect information on basic blocks:

- `bb_num`: Number of the current basic block.
- `shared_resources`: List of shared resources in the current block.
- `enable_disable_calls`: List of ISR activation/deactivation calls.
- `code_lines`: The lines of code of the current block.

- **line_number**: The current line number in the file.

The loop runs through each line of the file, removes leading and trailing spaces and increments the line number. If a new function is recognised (by the pattern `;; Function ... ()`), the current block is saved and the current function is updated.

If a new basic block number is detected (`<bb +>:`), the previous block is saved and the information for the new block is initialised.

This section searches each line for shared resources and ISR activation/deactivation calls. Resources and calls found are added to the corresponding list and the current line is added to `code_lines`.

At the end of the loop, the last basic block is saved if there is still an open block.

In a second pass through the lines, the successor blocks for each basic block are identified and linked accordingly. Finally, the function returns the `blocks` dictionary.

Algorithm 6: Track ISR Status

Data: blocks

Result: List indicating ISR status initialized to zero

Input: Dictionary of BasicBlocks

Output: List of zeros representing the status of each ISR

```
1 isr_count ← count unique ISRs in function names from blocks
2 return list initialized to zero of length isr_count
```

The function `track_isr_status` initialises the tracking of the ISR status. It returns a list of zeros whose length corresponds to the number of unique ISRs in the programme. This list is used to track the activation/deactivation status of each ISR.

Algorithm 7: Extract ISR Index

Data: function_name

Result: Zero-based index of the ISR or None if no ISR is found

Input: Function name as a string

Output: Integer index or None

```
1 match ← search for pattern 'isr[+]?(+)' in function_name
2 if match is found then
3   | return integer value of the first group in match minus one
4 else
5   | return None
6 end
```

The function `extract_isr_index` extracts the index of an ISR from a function name that follows the pattern `isr_+`. It returns the index of the ISR reduced by one to enable zero-based indexing.

Data Race Detection

Algorithm 8: Detect Data Races

Data: blocks
Result: List of potential data races identified
Input: Dictionary of BasicBlocks
Output: List of potential data races

```

1 potential_data_races ← empty list
2 resource_accesses ← initialize as a default dictionary to list
3 isr_enabling_map ← initialize as a default dictionary to set
4 foreach block in blocks do
5     foreach call, line_number in block.enable_disable_calls do
6         if call contains 'enable_isr' then
7             isr_idx_match ← search for pattern '+' in call
8             if isr_idx_match is found then
9                 enabled_isr_idx ← integer value of the first group in
10                     isr_idx_match minus one
11                 enabler_isr ← block.function_name
12                 isr_enabling_map[enabler_isr].add(enabled_isr_idx)
13             end
14         end
15 end
  
```

The first part of the function `detect_data_races` runs through all blocks and collects information about ISR activations. This information is stored in the `isr_enabling_map`, which tracks the ISR activation relationships.

Algorithm 9: Process Block

Data: block, current_isr_status**Result:** Updated ISR status and recorded resource accesses**Input:** A code block and the current ISR status as a list**Output:** Updated current ISR status and appended resource accesses to a global list

```

1 foreach line, line_number in block.code do
2   if line contains 'enable_isr' or 'disable_isr' then
3     isr_idx_match  $\leftarrow$  search for pattern '+' in line
4     if isr_idx_match is found then
5       isr_idx  $\leftarrow$  integer value of the first group in isr_idx_match minus one
6       if "disable_isr" is in line then
7         if  $0 \leq \text{isr\_idx} < \text{length of current\_isr\_status}$  then
8           current_isr_status[isr_idx]  $\leftarrow$  1
9         end
10      else
11        if  $0 \leq \text{isr\_idx} < \text{length of current\_isr\_status}$  then
12          current_isr_status[isr_idx]  $\leftarrow$  0
13        end
14      end
15    end
16  end
17  foreach resource_name, access_type, res_line_number in
    block.shared_resources do
18    if res_line_number == line_number then
19      resource_accesses[resource_name].append((block.function_name,
        block.number, access_type, line_number,
        current_isr_status.copy()))
20    end
21  end
22 end

```

The `process_block` function processes a single block and updates the current ISR status based on `enable_isr` and `disable_isr` calls. Accesses to shared resources are stored in `resource_accesses`.

Algorithm 10: Depth-First Search (DFS) for Blocks

Data: block, visited_blocks, current_isr_status, path**Result:** Recursive processing of blocks and their successors**Input:** A block, set of visited blocks, current ISR status, and path tracing the blocks visited**Output:** Processed blocks with updated ISR statuses and path

```

1 if (block.function_name, block.number) is in visited_blocks then
2   | return
3 end

4 visited_blocks.add((block.function_name, block.number))
5 path.append((block.function_name, block.number))
6 process_block(block, current_isr_status)
7 if block has successors then
8   | foreach successor in block.successors do
9     |   dfs(successor, set(visited_blocks), current_isr_status.copy(), path.copy())
10  | end
11 end

```

The function `dfs` (Depth-First Search) runs through the control flow graph and tracks the ISR status. Each block is processed and the ISR status and accesses to resources are updated.

Algorithm 11: Initial Processing and Recursive Depth-First Search

Data: blocks (a dictionary of blocks indexed by function name and block number)**Result:** Executes process block and DFS on selected blocks and their successors

```

1 foreach (func_name, bb_num) and block in blocks do
2   | if bb_num = 2 then
3     |   initial_isr_status ← track_isr_status(blocks)
4     |   process_block(block, initial_isr_status)
5     |   foreach successor in block.successors do
6       |     dfs(successor, {}, initial_isr_status.copy(), [(func_name, bb_num)])
7     |   end
8   | end
9 end

```

This section initialises the depth search for basic blocks with the number 2 and starts processing and running through the successor blocks.

Algorithm 12: Check for Data Races

Result: Update global list of potential data races based on access patterns

Output: Updated list of potential data races with detailed conflict info

```

1 foreach resource, accesses in resource_accesses do
2   for  $i \leftarrow 0$  to  $\text{length}(\text{accesses})-1$  do
3      $\text{func1}, \text{bb\_num1}, \text{access\_type1}, \text{line\_number1}, \text{isr\_status1} \leftarrow \text{accesses}[i]$ 
4     for  $j \leftarrow i+1$  to  $\text{length}(\text{accesses})-1$  do
5        $\text{func2}, \text{bb\_num2}, \text{access\_type2}, \text{line\_number2}, \text{isr\_status2} \leftarrow$ 
6          $\text{accesses}[j]$ 
7       if  $\text{func1} \neq \text{func2}$  and  $(\text{access\_type1} = \text{"write"} \text{ or } \text{access\_type2} =$ 
8          $\text{"write"})$  then
9          $\text{potential\_data\_races.append}((\text{resource}, (\text{func1}, \text{bb\_num1},$ 
10           $\text{access\_type1}, \text{line\_number1}, \text{isr\_status1}), (\text{func2}, \text{bb\_num2},$ 
11            $\text{access\_type2}, \text{line\_number2}, \text{isr\_status2})))$ 
8       end
9     end
10   end
11 end

```

The function `check_for_data_races` identifies potential data races by comparing pairs of resource accesses. If two accesses to the same resource originate from different functions and at least one of them is a write access, a potential data race is identified.

Calling `check_for_data_races` performs a check of the potential data races.

Filter Possible Data Races

Algorithm 13: Filter Data Races

Data: *potential_data_races*

Result: Filtered list of data races

```

1 Function filter_data_races(potential_data_races):
2   filtered_data_races  $\leftarrow$  empty list
3   foreach resource, access1, access2 in potential_data_races do
4     Decompose access1 into func1, bb_num1, access_type1, line_number1,
       isr_status1
5     Decompose access2 into func2, bb_num2, access_type2, line_number2,
       isr_status2
6     relevant_isr_disabled1  $\leftarrow$  is_isr_disabled(isr_status1, func2) and not
       is_isr_enabled_by_another(isr_status1, func2)
7     relevant_isr_disabled2  $\leftarrow$  is_isr_disabled(isr_status2, func1) and not
       is_isr_enabled_by_another(isr_status2, func1)
8     if not (relevant_isr_disabled1 or relevant_isr_disabled2) then
9       | filtered_data_races.append((resource, access1, access2))
10    end
11  end
12  return filtered_data_races
13 end
14 Function is_isr_disabled(isr_status, func_name):
15   isr_idx  $\leftarrow$  extract_isr_index(func_name)
16   if isr_idx is not None and isr_idx < length of isr_status then
17     | return isr_status[isr_idx] == 1
18     | // 1 indicates disabled
19   end
20   return False
21 Function is_isr_enabled_by_another(isr_status, func_name):
22   isr_idx  $\leftarrow$  extract_isr_index(func_name)
23   if isr_idx is not None then
24     | foreach enabler_isr, enabled_isrs in isr_enabling_map do
25       | enabler_idx  $\leftarrow$  extract_isr_index(enabler_isr)
26       | if enabler_idx is not None and not is_isr_disabled(isr_status,
27         | enabler_isr) then
28         | | if isr_idx in enabled_isrs then
29         | | | return True
30         | | end
31       | end
32     | end
33     return False
34 end

```

The function `filter_data_races` filters false positives from the list of potential data races. It checks the ISR status during resource accesses and only confirms data races if relevant ISRs were not deactivated at access time or activated by another function. At the end, `detect_data_races` returns the filtered data races.

4 Evaluation

1,5 Woche

5 Conclusion

Indroduction+Conclusion und Allgemeine Überarbeitung 0,5 Woche 1 Wochen Korrekturlesen und Einarbeitung =9 Wochen bei Vollarbeitszeit an BE

Bibliography

Lightweight Data Race Detection for Production by Swarnendu Biswas, Man Cao, Minjia Zhang, Michael D. Bond, Benjamin P. Wood

A Deployable Sampling Strategy for Data Race Detection by Yan Cai¹, Jian Zhang, Lingwei Cao, and Jian Liu

Bibliography

- [1] Wang, Y., Gao, F., Wang, L., Yu, T., Zhao, J., & Li, X. (2020). Automatic Detection, Validation, and Repair of Race Conditions in Interrupt-Driven Embedded Software. *IEEE Transactions on Software Engineering*.
- [2] Engler, D., & Ashcraft, K. (2003). RacerX: Effective, Static Detection of Race Conditions and Deadlocks. *ACM SIGOPS Operating Systems Review*.
- [3] Nikita Chopra, Rekha Pai, and Deepak D'Souza (2019). Data Races and Static Analysis for Interrupt-Driven Kernels
- [4] Flanagan, C., & Freund, S. N. (2009). FastTrack: Efficient and Precise Dynamic Race Detection. *ACM SIGPLAN Notices*.
- [5] R. Chen, Xiangying Guo, Y. Duan, B. Gu, Mengfei Yang (2011). Static Data Race Detection for Interrupt-Driven Embedded Software.
- [6] Muchnick, S. S. (1997). *Advanced Compiler Design and Implementation*. Morgan Kaufmann.
- [7] Adve, S. V., & Gharachorloo, K. (1996). Shared Memory Consistency Models: A Tutorial. *IEEE Computer*.
- [8] Herlihy, M., & Shavit, N. (2008). *The Art of Multiprocessor Programming*. Morgan Kaufmann.

Attachments

```

1  class BasicBlock:
2  def __init__(function_name, number, shared_resources=[], successors
    =[], enable_disable_calls=[], code=[]):
3      self.function_name = function_name
4      self.number = number
5      self.shared_resources = shared_resources
6      self.successors = successors
7      self.enable_disable_calls = enable_disable_calls
8      self.code = code
9
10 def __repr__():
11 return ("BasicBlock(function_name={}, number={}, shared_resources={},
    "
12 "successors={}, enable_disable_calls={}, code={})".format(
13 self.function_name, self.number, self.shared_resources,
14 [succ.number for succ in self.successors], self.enable_disable_calls,
15 ' '.join(self.code)))
16
17 def parse_basic_blocks(file_path, shared_resource_names):
18     blocks = {}
19     current_function = None
20
21     with open(file_path, 'r') as file:
22         lines = file.readlines()
23
24     bb_num = None
25     shared_resources = []
26     enable_disable_calls = []
27     code_lines = []
28     line_number = 0
29
30     for line in lines:
31         line = line.strip()
32         line_number += 1
33
34         func_match = re.match(r';; Function (.+?) \(', line)
35         if func_match:
36             if bb_num is not None and current_function is not None:
37                 blocks[(current_function, bb_num)] = BasicBlock(
38                     current_function, bb_num, shared_resources, [], enable_disable_calls,
39                     code_lines)
40             current_function = func_match.group(1)
41             bb_num = None
42             continue
43
44         bb_match = re.match(r'<bb (\d+)>:', line)
45         if bb_match:
46             if bb_num is not None and current_function is not None:
47                 blocks[(current_function, bb_num)] = BasicBlock(
48                     current_function, bb_num, shared_resources, [], enable_disable_calls,
49                     code_lines)

```

```

48     bb_num = int(bb_match.group(1))
49     shared_resources = []
50     enable_disable_calls = []
51     code_lines = []
52
53     for resource_name in shared_resource_names:
54         if re.search(fr'\b{resource_name}\b', line):
55             if re.search(fr'\b{resource_name}\b\s*=', line):
56                 shared_resources.append((resource_name, 'write', line_number))
57             else:
58                 shared_resources.append((resource_name, 'read', line_number))
59
60         if 'enable_isr' in line or 'disable_isr' in line:
61             enable_disable_calls.append((line.strip(), line_number))
62
63     code_lines.append((line, line_number))
64
65     if bb_num is not None and current_function is not None:
66         blocks[(current_function, bb_num)] = BasicBlock(
67             current_function, bb_num, shared_resources, [], enable_disable_calls,
68             code_lines)
69
70     current_function = None
71     bb_num = None
72     for line in lines:
73         line = line.strip()
74
75         func_match = re.match(r';; Function (.+?) \(', line)
76         if func_match:
77             current_function = func_match.group(1)
78             bb_num = None
79             continue
80
81         if 'succs' in line:
82             succ_match = re.match(r';; (\d+) succs \{(.+?)\}', line)
83             if succ_match:
84                 bb_num = int(succ_match.group(1))
85                 succ_list = [int(succ.strip()) for succ in succ_match.group(2).split(
86                     ' ')]
87                 if (current_function, bb_num) in blocks:
88                     blocks[(current_function, bb_num)].successors = [
89                         blocks[(current_function, succ)] for succ in succ_list if (
90                             current_function, succ) in blocks]
91
92     return blocks
93
94     def track_isr_status(blocks):
95         isr_count = len(set(block.function_name for block in blocks.values()
96             if re.search(r'isr[_]?(\d+)', block.function_name)))
97         return [0] * isr_count
98
99     def extract_isr_index(function_name):
100         match = re.search(r'isr[_]?(\d+)', function_name)
101         if match:

```

```

98     return int(match.group(1)) - 1
99     return None
100
101     def detect_data_races(blocks):
102         potential_data_races = []
103         resource_accesses = defaultdict(list)
104         isr_enabling_map = defaultdict(set)
105
106         for block in blocks.values():
107             for call, line_number in block.enable_disable_calls:
108                 if 'enable_isr' in call:
109                     isr_idx_match = re.search(r'\((\d+)\)', call)
110                     if isr_idx_match:
111                         enabled_isr_idx = int(isr_idx_match.group(1)) - 1
112                         enabler_isr = block.function_name
113                         isr_enabling_map[enabler_isr].add(enabled_isr_idx)
114
115         def process_block(block, current_isr_status):
116             for line, line_number in block.code:
117                 if 'enable_isr' in line or 'disable_isr' in line:
118                     isr_idx_match = re.search(r'\((\d+)\)', line)
119                     if isr_idx_match:
120                         isr_idx = int(isr_idx_match.group(1)) - 1
121                         if "disable_isr" in line:
122                             if 0 <= isr_idx < len(current_isr_status):
123                                 current_isr_status[isr_idx] = 1
124                         elif "enable_isr" in line:
125                             if 0 <= isr_idx < len(current_isr_status):
126                                 current_isr_status[isr_idx] = 0
127
128
129             for resource_name, access_type, res_line_number in block.
130                 shared_resources:
131                 if res_line_number == line_number:
132                     resource_accesses[resource_name].append((block.function_name, block.
133                         number, access_type, line_number, current_isr_status.copy()))
134
135         def dfs(block, visited_blocks, current_isr_status, path):
136             if (block.function_name, block.number) in visited_blocks:
137                 return
138             visited_blocks.add((block.function_name, block.number))
139             path.append((block.function_name, block.number))
140
141             process_block(block, current_isr_status)
142
143             if not block.successors:
144                 pass
145             else:
146                 for successor in block.successors:
147                     dfs(successor, set(visited_blocks), current_isr_status.copy(), path.
148                         copy())
149
150         for (func_name, bb_num), block in blocks.items():

```

```

149     if bb_num == 2:
150         initial_isr_status = track_isr_status(blocks).copy()
151         process_block(block, initial_isr_status)
152         for successor in block.successors:
153             dfs(successor, set(), initial_isr_status.copy(), [(func_name, bb_num)
154                 ])
155
156     def check_for_data_races():
157         for resource, accesses in resource_accesses.items():
158             for i, (func1, bb_num1, access_type1, line_number1, isr_status1) in
159                 enumerate(accesses):
160                 for j, (func2, bb_num2, access_type2, line_number2, isr_status2) in
161                     enumerate(accesses):
162                     if i >= j:
163                         continue
164                     if func1 != func2 and (access_type1 == "write" or access_type2 == "
165                         write"):
166                         potential_data_races.append((resource, (func1, bb_num1, access_type1,
167                             line_number1, isr_status1),
168                             (func2, bb_num2, access_type2, line_number2, isr_status2)))
169
170     check_for_data_races()
171
172     def filter_data_races(potential_data_races):
173         filtered_data_races = []
174         for resource, access1, access2 in potential_data_races:
175             func1, bb_num1, access_type1, line_number1, isr_status1 = access1
176             func2, bb_num2, access_type2, line_number2, isr_status2 = access2
177
178     def is_isr_disabled(isr_status, func_name):
179         isr_idx = extract_isr_index(func_name)
180         if isr_idx is not None and isr_idx < len(isr_status):
181             return isr_status[isr_idx] == 1
182         return False
183
184     def is_isr_enabled_by_another(isr_status, func_name):
185         isr_idx = extract_isr_index(func_name)
186         if isr_idx is not None:
187             for enabler_isr, enabled_isrs in isr_enabling_map.items():
188                 enabler_idx = extract_isr_index(enabler_isr)
189                 if enabler_idx is not None and not is_isr_disabled(isr_status,
190                     enabler_isr):
191                     if isr_idx in enabled_isrs:
192                         return True
193             return False
194
195     relevant_isr_disabled1 = is_isr_disabled(isr_status1, func2) and not
196         is_isr_enabled_by_another(isr_status1, func2)
197     relevant_isr_disabled2 = is_isr_disabled(isr_status2, func1) and not
198         is_isr_enabled_by_another(isr_status2, func1)
199
200     if not (relevant_isr_disabled1 or relevant_isr_disabled2):

```

```
195     filtered_data_races.append((resource, access1, access2))
196
197     return filtered_data_races
198
199     filtered_data_races = filter_data_races(potential_data_races)
200
201     return filtered_data_races
202
203
204     shared_resource_input = input("Enter the names of shared resources,
205                                   separated by commas: ")
206     shared_resource_names = [name.strip() for name in
207                             shared_resource_input.split(',')]
208
209     file_path = input("Enter the file path: ").strip()
210     blocks = parse_basic_blocks(file_path, shared_resource_names)
211
212     data_races = detect_data_races(blocks)
213
214     print("Detected Data Races:")
215     for resource, access1, access2 in data_races:
216         print(f"Resource: {resource}")
217         print(f"  Access 1: Function {access1[0]} (BB {access1[1]}), {access1[2]}, Line {access1[3]}, ISR Status: {access1[4]}")
218         print(f"  Access 2: Function {access2[0]} (BB {access2[1]}), {access2[2]}, Line {access2[3]}, ISR Status: {access2[4]}")
219         print()
```