

Robin Willenbrock

# **Static Detection of Data Races in Interrupt-Driven Software Using Reduced Inter-Procedural Control Flow Graphs**

July 2, 2024

---

supervised by:

Prof. Dr. Sibylle Schupp  
Ulrike Engeln

---

Hamburg University of Technology (TUHH)  
*Technische Universität Hamburg*  
Institute for Software Systems  
21073 Hamburg



# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Background</b>	<b>3</b>
2.1	Interrupt-Driven Systems . . . . .	3
2.2	Shared Resources . . . . .	4
2.3	Reduced Inter-Procedural Control Flow Graphs . . . . .	4
2.3.1	Reduction of Control Flow Graphs . . . . .	5
2.3.2	Depth-First Search . . . . .	6
2.4	Data Races . . . . .	7
2.4.1	Detection Techniques . . . . .	8
2.4.2	Strategies for Preventing Data Races . . . . .	9
2.5	Static Detection of Data Races in Interrupt-Driven Systems . . . . .	10
<b>3</b>	<b>Implementation</b>	<b>13</b>
<b>4</b>	<b>Evaluation</b>	<b>27</b>
<b>5</b>	<b>Conclusion</b>	<b>29</b>
	<b>Attachments</b>	<b>33</b>



# List of Figures

2.1	Flowchart of the Interrupt-Driven System . . . . .	3
2.2	Example of an Inter-Procedural Control Flow Graph . . . . .	5
2.3	Example of a Reduced Inter-Procedural Control Flow Graph . . . . .	6
2.4	Example Implementation of a Depth-First Search . . . . .	7
2.5	Simple Example of a Data Race . . . . .	8
2.6	Example of a Data Race with Enable/Disable Interrupt Service Routine (ISR) Calls . . . . .	10
2.7	Static Race Detection Approach by [1] . . . . .	11
3.1	Algorithm: Class BasicBlock . . . . .	13
3.2	UML: Class BasicBlock . . . . .	14
3.3	Algorithm: Initialization and ISR Enabling Map . . . . .	15
3.4	Algorithm: Process Block . . . . .	16
3.5	Algorithm: Depth-First Search (DFS) and Initialization . . . . .	17
3.6	Algorithm: Merge ISR Statuses . . . . .	18
3.7	Algorithm: Check for Data Races . . . . .	19
3.8	Algorithm: Filter Data Races . . . . .	20
3.9	Algorithm: Is ISR Disabled . . . . .	21
3.10	Algorithm: Is ISR Enabled by Another . . . . .	21
3.11	Algorithm: Parse Basic Blocks . . . . .	23
3.12	Algorithm: Determine Priority . . . . .	24
3.13	Algorithm: Track ISR Status . . . . .	24
3.14	Algorithm: Track ISR Status . . . . .	25
3.15	Algorithm: Propagate Function Calls . . . . .	25



# Abbreviations

**ISR** Interrupt Service Routine

**CFG** Control Flow Graph

**ICFG** Inter-Procedural Control Flow Graph

**RICFG** Reduced Inter-Procedural Control Flow Graph

**GCC** GNU Compiler Collection

**DFS** Depth-First Search

**BB** BasicBlock





# 1 Introduction



## 2 Background

In this chapter, I am going to provide a brief overview of all the necessary background information needed to understand static data races in interrupt-driven software using Reduced Inter-Procedural Control Flow Graphs (RICFGs). This information includes basics about interrupt-driven systems, shared resources, RICFGs, and data races as a whole.

### 2.1 Interrupt-Driven Systems

An interrupt-driven system is an architecture where the flow of execution is changed by unpredictable events in the system, also known as interrupts. Interrupts can be caused by hardware devices, software conditions, or external signals, forcing the processor to suspend the current task to execute an interrupt handler or Interrupt Service Routine (ISR). Interrupt-driven systems are used in real-time operating systems, embedded systems, and generally in systems where timely responses are necessary [1].

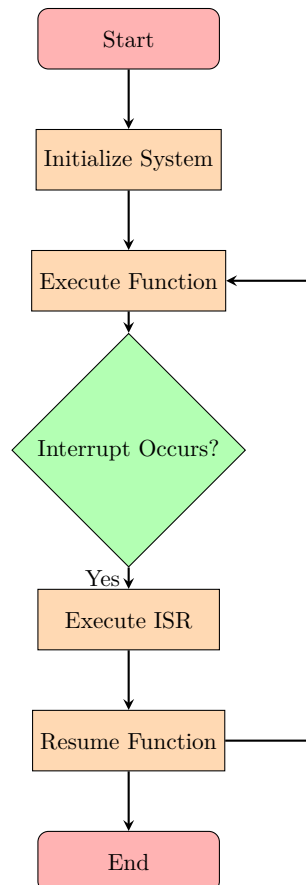


Figure 2.1: Flowchart of the Interrupt-Driven System

In Figure 2.1, a basic execution flow of a simple interrupt-driven system is displayed. The system executes a function as long as no interrupt occurs. When an interrupt occurs, it switches to the ISR, executes it, and then resumes the function executed before the interrupt happened.

The management of interrupts to maintain the fast responsiveness of the system is the most challenging part of an interrupt-driven system. Interrupts occur in unpredictable ways, so you have to consider every possible execution flow. To ensure the execution of critical interrupts, interrupts are often prioritized, so higher-priority events can interrupt lower ones and be handled immediately. When handling an interrupt, the current state of the processor is saved, and the context is switched to the ISR [1].

The unpredictability and asynchronous nature of interrupts present many challenges in designing and implementing an interrupt-driven system. One of the biggest challenges is the correct handling of high-priority interrupts without delaying them substantially, which requires a sophisticated scheduling and prioritization mechanism. The execution of the main program and ISR needs to be handled properly to ensure data integrity.

## 2.2 Shared Resources

Analyzing the management of shared resources is a large part of data race analysis, which is further explained later. The following is a short introduction to shared resources to better understand them in the context of data races.

Shared resources, often referred to as shared memory or shared variables, are data that can be accessed simultaneously by multiple threads or processes. Proper management of these resources is crucial because improper handling can lead to issues like data races, deadlocks, and other synchronization problems. In interrupt-driven systems, shared resources often involve variables or data structures that are accessed by both the main program and ISRs. Proper management of shared resources is critical to ensuring data consistency and avoiding conflicts [8].

Proper management of shared resources involves the use of synchronization mechanisms to coordinate access and ensure data consistency. Mutexes, semaphores, and condition variables are common tools used to control access to shared resources. Mutexes provide mutual exclusion, ensuring that only one thread can access the resource at a time. Semaphores can limit the number of threads accessing the resource simultaneously. Condition variables allow threads to wait for certain conditions to be met before proceeding, facilitating complex synchronization scenarios [8]. In interrupt-driven software, the synchronization of shared resources often involves disabling and enabling interrupts [3].

## 2.3 Reduced Inter-Procedural Control Flow Graphs

Control Flow Graphs (CFGs) are representations of all possible paths through a program or function during its execution. An Inter-Procedural Control Flow Graphs (ICFGs) adds possible edges between multiple programs or functions to also show possible control

flows between those. A RICFG is an optimized version of the ICFG that simplifies the graph to only the necessary information needed for the analysis [2].

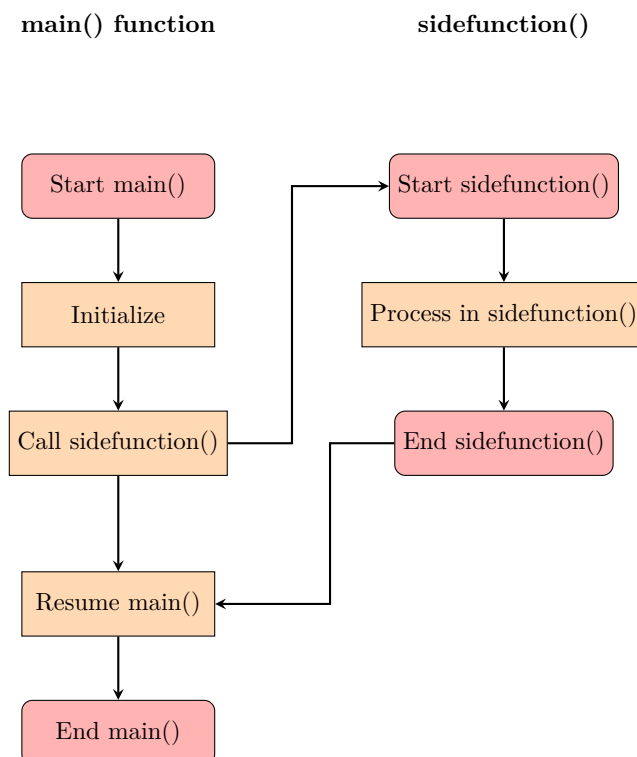


Figure 2.2: Example of an Inter-Procedural Control Flow Graph

In Figure 2.2, a simple ICFG is shown. There are two separate linear control flow graphs where the main function calls the side function in its execution. To interpret the flow of the program correctly, you need to consider the execution of `sidefunction()` and where it's called. The ICFG combines the two separate CFGs to ensure correct analysis.

### 2.3.1 Reduction of Control Flow Graphs

There are multiple techniques to reduce the graph, such as node merging, edge contraction, and the elimination of non-important nodes, without losing any information required for the analysis and reducing the complexity of the RICFG. The reduction of the ICFG makes the analysis of large and complex software a lot more efficient. By minimizing the amount of data while retaining enough detail, RICFGs are great for static analysis of data races [1].

Node merging involves combining nodes that represent redundant control flow paths to reduce the number of nodes in the graph. Edge contraction simplifies the graph by reducing the number of edges between nodes. It collapses edges that do not significantly affect the control flow of the graph [6]. The elimination of nodes is the main tool used

in this work to reduce the CFG. Eliminating nodes that do not carry any essential information for the applied data analysis significantly reduces the amount of data the algorithm has to analyze. Overall, these techniques enhance the scalability of static analysis and make it more practical to analyze more complex data [1].

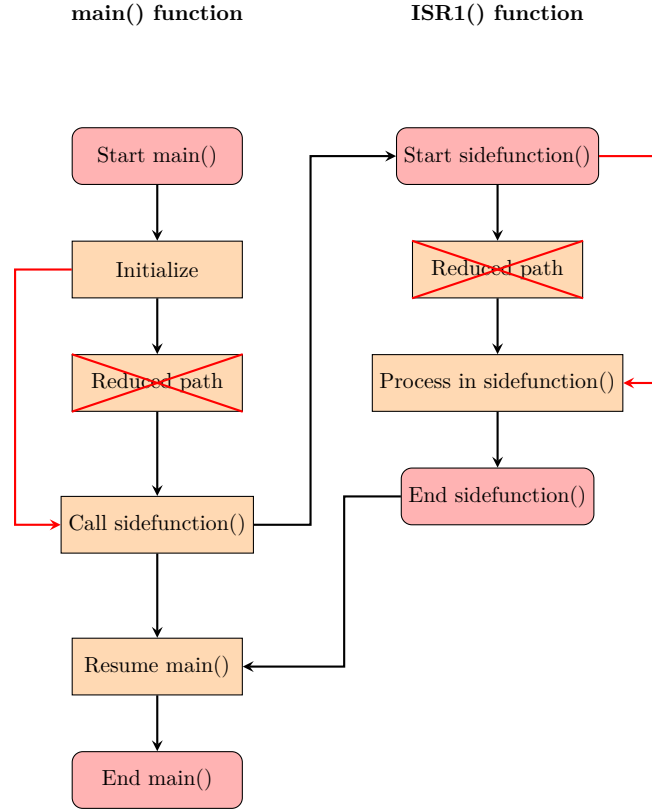


Figure 2.3: Example of a Reduced Inter-Procedural Control Flow Graph

Figure 2.3 shows an example of a simple reduction by eliminating nodes that do not carry any important information for the analysis the RICFG is used for.

### 2.3.2 Depth-First Search

DFS is an algorithm used to traverse a graph systematically. It begins at a source node and extends its exploration through the connected nodes as far as possible before it backtracks. In an algorithm this can be implemented using a recursive approach. The basic idea is to mark a node when its first discovered and then explore all the adjacent nodes that are not visited before.

---

```
1 Function dfs(node, visited):  
2   visited[node] = True;  
3   for neighbor in graph[node] do  
4     if not visited[neighbor] then  
5       dfs(neighbor, visited);  
6     end  
7   end
```

---

Figure 2.4: Example Implementation of a Depth-First Search

In Figure 2.4 a simple example of a DFS is shown. The `dfs` function gets called with the root node of a tree and performs the DFS starting from the given node. It recursively calls itself with a new node called `neighbor` and repeats this until all the reachable nodes are marked as visited [9].

## 2.4 Data Races

A data race occurs when two or more functions or threads access a shared resource concurrently, without being ordered by a happens-before relationship, and one of those accesses is a write operation [5]. This can lead to unpredictable behavior and errors in the system, which makes the detection of data races a critical aspect of concurrent programs. Without proper synchronization, a system with multiple threads or functions that use shared data will lead to data races. The outcome of a program with data races is non-deterministic [5]. The order of execution of operations can vary, which may result in the generation of bugs that are not reproducible or difficult to reproduce.

---

**Algorithm 1:** Data Race Example

---

```

Data: long shared1;
1 Function main ():
2   Variables:
3   unsigned char tmp;
4   Code:
5   tmp ← shared1;
6 Function isr1 ():
7   Code:
8   idlerun();
9   shared1 ← 1;
10  idlerun();

```

---

Figure 2.5: Simple Example of a Data Race

In Figure 2.4, an example of a simple data race is shown. A global variable `shared1` is initiated and accessed in two different functions, `main()` and `isr1()`. Since there are no synchronization tools used and the operation in `isr1` is a write, there is a data race between line 5 and line 9.

### 2.4.1 Detection Techniques

Data race detection can be approached by two different analytical methods. Each of these methods provides benefits and challenges.

#### Static Data Race Detection [1]

##### Advantages:

- **Comprehensiveness:** Static analysis inspects the code without executing the program by analyzing every possible execution path and interaction that could lead to data races.
- **Early Detection:** Since static analysis does not require execution, it can analyze the code in the development phase, allowing the developer to find issues without deployment.

##### Disadvantages:

- **False Positives/Negatives:** Static analysis reports all data races that fall under certain conditions. Some of these data races could be very unlikely or even impossible at runtime. On the other hand, due to the approximations and assumptions necessary for tractability, it may miss some races.



- **Complexity in Handling Dynamic Behavior:** Dynamic behaviors such as pointers or recursion can be challenging to analyze for static approaches, leading to incomplete or inaccurate results.

### Dynamic Data Race Detection [4]

#### Advantages:

- **Precision:** Dynamic analysis tools monitor the actual execution of a program, identifying data races in real-time, which reduces the number of false positives.
- **Context-Sensitive Detection:** By analyzing the actual runtime behavior, dynamic analysis can understand the context of operations, leading to more accurate detection.

#### Disadvantages:

- **Performance Overhead:** The analysis at runtime can slow down the application significantly.
- **Coverage:** The effectiveness is heavily dependent on the execution path triggered during the tests. If certain parts of the program are not passed through in the execution run, they are not analyzed.

Both static and dynamic analyses are crucial for a complete analysis of code. They complement each other's limitations. A combination of both is the best approach to detecting data races most reliably. However, in this work, I am going to focus on the static analysis of data races.

### 2.4.2 Strategies for Preventing Data Races

Preventing data races requires careful design and implementation of concurrent programs. Effective strategies for general prevention of data races are synchronization mechanisms such as mutexes, semaphores, and condition variables, which control access to shared data. These mechanisms ensure that only one thread can access the shared resource at a time [8]. Since I am focusing on data races in interrupt-driven systems, the main tool to prevent data races is to disable ISRs, which access shared resources in critical areas.

---

**Algorithm 2:** Enable/Disable ISR Call Example
 

---

```

Data: long shared1
1 Function main():
2   Variables:
3   unsigned char tmp;
4   Code:
5   disable_isr(1);
6   tmp ← shared1;
7   enable_isr(1);
8 Function isr1():
9   Code:
10  idlerun();
11  shared1 ← 1;
12  idlerun();
13 Function isr2():
14  Code:
15  idlerun();
16  int variable1 = 1;
17  idlerun();

```

---

Figure 2.6: Example of a Data Race with Enable/Disable ISR Calls

Figure 2.5 is an example of a disable ISR call that leads to the safe access of the shared data. The main function and *isr1* both access the shared resource *shared1*. Since the read operation in line 6 of the main function is safely accessed by disabling *isr1* in line 5 and enabling it in line 7, a possible data race is prevented.

## 2.5 Static Detection of Data Races in Interrupt-Driven Systems

The asynchronous nature and concurrent execution of ISRs and the main function introduce significant challenges for data consistency and detecting data races in interrupt-driven systems. Static data race analysis, especially those using RICFGs, is a promising approach to identifying data races without the need for extensive testing and runtime monitoring as in dynamic approaches [1].

The static approach involves the construction of an RICFG for the program, which includes both the main code and ISRs, and capturing the control flow and potential interaction between them. Analyzing the RICFG shows paths where shared resources are accessed concurrently without proper synchronization and indicates potential data races. Integrating the static analysis tool with the development process enables continuous detection of data races during software development, improving the reliability and

correctness of interrupt-driven systems [1].

The methodology for static data race detection in interrupt-driven systems involves the following key steps. First, the RICFGs are constructed for the entire program, including the main code and the ISRs. This involves analyzing the control flow and identifying interactions between the main program and ISRs. Next, the RICFGs are analyzed to find potential data races, focusing on paths where concurrent access to shared data is done without proper synchronization. Finally, the developer can use the analysis results to address identified data races early in the development process [1].

---

**Algorithm 3:** Static Race Detection
 

---

**Input:** RICFGs of P

**Output:** potential racing pairs (PR)

```

1 for each  $\langle G_i; G_j \rangle$  in RICFGs do
2   for each  $sv_i \in G_i$  do
3     for each  $sv_j \in G_j$  do
4       if  $sv_i.V == sv_j.V$  and  $(sv_i.A == W$  or  $sv_j.A == W)$  and
5          $G_i.pri < G_j.pri$  and  $INTB.get(sv_i).contains(G_j)$  then
           $PR = PR \cup \{\langle sv_i, sv_j \rangle\};$ 

```

---

Figure 2.7: Static Race Detection Approach by [1]

The approach by Wang et al. shows a computation of potential data races using RICFGs. By running a depth-first search on the RICFGs, it finds the interrupt status of every instruction. If there is a shared resource in both of the analyzed RICFGs, at least one of them is a write operation, and the two functions differ in their priority. While the interrupt in this pair is enabled, the two accesses are a potential data race [1].

In the following, I am going to introduce you to the implementation of my static analysis program based on the static race detection approach of Wang et al.



## 3 Implementation

In the following, I will provide an in-depth explanation of my implementation. For the generation of the input, I used GNU Compiler Collection (GCC). The command `gcc -fdump-tree-cfg` provides a `cfg`-file with all the important information for the intended data race analysis. I have split the explanation of the implementation into the initialization of the basic block class, the parsing of the input, the actual data race analysis, and the filtering of false positives found in data race analysis.

### Class BasicBlock

---

**Algorithm 4:** BasicBlock Constructor

---

**Data:** `function_name`, `number`, `priority`, `shared_resources`, `successors`,  
`enable_disable_calls`, `function_calls`

**Result:** A BasicBlock object

**Input:** `function_name`, `number`, `priority`, `shared_resources` (optional), `successors`  
 (optional), `enable_disable_calls` (optional), `function_calls` (optional)

**Output:** A BasicBlock object

```

1 Function BasicBlock(function_name, number, priority, shared_resources,
  successors, enable_disable_calls, function_calls):
2   self.function_name ← function_name
3   self.number ← number
4   self.priority ← priority
5   self.shared_resources ← shared_resources if shared_resources else empty list
6   self.successors ← successors if successors else empty list
7   self.enable_disable_calls ← enable_disable_calls if enable_disable_calls
  else empty list
8   self.function_calls ← function_calls if function_calls else empty list

```

---

Figure 3.1: Algorithm: Class BasicBlock

The class `BasicBlock` displays all the information necessary for the data race analysis found in the input. This information includes the following attributes:

- **function\_name:** The function name to which the basic block belongs.
- **number:** The number of the basic block.
- **shared\_resources:** All accesses of shared resources within the BasicBlock (BB). The access type (read/write) and the line number of such calls are saved.

- **successors**: A list of all the successors of each basic block. Important for building all possible paths through the CFG.
- **enable\_disable\_calls**: All calls that disable or enable an ISR within this basic block and also the corresponding line number of those calls to ensure the correct order.
- **function\_calls**: The functions that are called within a BB.

Resulting in the following UML class diagramm:

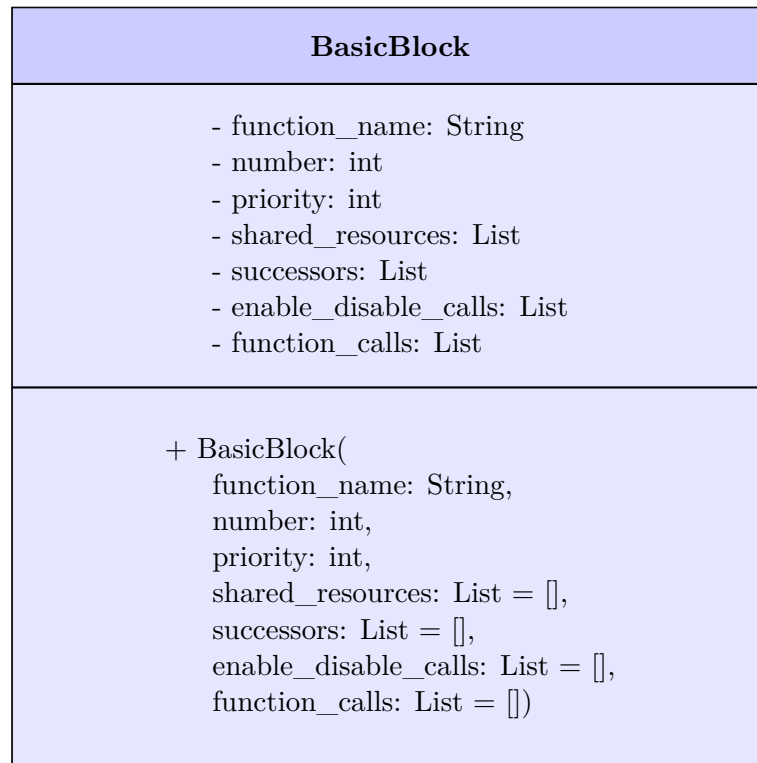


Figure 3.2: UML: Class BasicBlock

## Data Race Detection

The following functions are used to determine all possible data races, which are filtered later in the code. The intention is to find all possible data races to minimize the number of false negatives. Since false positives can be evaluated later by interpreting the output.

---

**Algorithm 5:** Initialization and ISR Enabling Map

---

**Data:** blocks  
**Result:** List of potential data races identified  
**Input:** Dictionary of BasicBlocks  
**Output:** List of potential data races

```

1 potential_data_races ← empty list
2 resource_accesses ← initialize as a default dictionary to list
3 isr_enabling_map ← initialize as a default dictionary to set
4 foreach block in blocks do
5   foreach call, line_number in block.enable_disable_calls do
6     if call contains 'enable_isr' then
7       isr_idx_match ← search for pattern '+' in call
8       if isr_idx_match is found then
9         enabled_isr_idx ← integer value of the first group in
          isr_idx_match minus one
10        enabler_isr ← block.function_name
11        isr_enabling_map[enabler_isr].add(enabled_isr_idx)
12      end
13    end
14  end
15 end

```

---

Figure 3.3: Algorithm: Initialization and ISR Enabling Map

The first part of the function `detect_data_races` takes a list of all basic block items as input. It also initializes the empty list of `potential_data_races`, a dictionary for `resource_access`, and a dictionary for the `isr_enabling_map`. `Potential_data_races` and `resource_access` are used later in the code. The main loop of the function iterates through every item in `blocks` and finds basic blocks with `enable_disable_calls`. If there is an enable call in a block item, the index of the enabled ISR is read, and the basic block is added to the `isr_enabling_map` with the information on which ISR it enables.

**Algorithm 6:** Process Block

---

**Data:** block, current\_isr\_status  
**Result:** Updated ISR status and recorded resource accesses  
**Input:** A code block and the current ISR status as a list  
**Output:** Updated current ISR status and appended resource accesses to a global list

```

1 foreach line, line_number in block.code do
2   if line contains 'enable_isr' or 'disable_isr' then
3     isr_idx_match  $\leftarrow$  search for pattern '+' in line
4     if isr_idx_match is found then
5       isr_idx  $\leftarrow$  integer value of the first group in isr_idx_match minus one
6       if "disable_isr" is in line then
7         if  $0 \leq \text{isr\_idx} < \text{length of current\_isr\_status}$  then
8           current_isr_status[isr_idx]  $\leftarrow$  1
9         end
10      else
11        if  $0 \leq \text{isr\_idx} < \text{length of current\_isr\_status}$  then
12          current_isr_status[isr_idx]  $\leftarrow$  0
13        end
14      end
15    end
16  end
17  foreach resource_name, access_type, res_line_number in
    block.shared_resources do
18    if res_line_number == line_number then
19      resource_accesses[resource_name].append((block.function_name,
        block.number, access_type, line_number,
        current_isr_status.copy()))
20    end
21  end
22 end

```

---

Figure 3.4: Algorithm: Process Block

This part of the code runs through each line of a basic block to find the current ISR status while resources are accessed. When a resource is found, all the information is added to the **resource\_accesses** dictionary, which includes the function name and the block number of the current basic block, as well as the access type, the line number, and the ISR status of the access. All this information is used later for the detection of data races.



---

**Algorithm 7:** Depth-First Search (DFS) and Initialization

---

**Data:** blocks**Result:** Updated block ISR statuses and processed blocks**Input:** Dictionary of basic blocks**Output:** Updated block ISR statuses

```

1 Function dfs(block, visited_blocks, current_isr_status, path):
2   if (block.function_name, block.number) in visited_blocks then
3     block_isr_statuses[(block.function_name, block.number)] ←
4       merge_isr_statuses(block_isr_statuses[(block.function_name,
5         block.number)], current_isr_status)
6   return
7 end
8 visited_blocks.add((block.function_name, block.number))
9 path.append((block.function_name, block.number))
10 block_isr_statuses[(block.function_name, block.number)] ←
11   current_isr_status.copy()
12 process_block(block, current_isr_status)
13 if not block.successors then
14   return
15 end
16 else
17   for successor in block.successors do
18     dfs(successor, set(visited_blocks), current_isr_status.copy(),
19       path.copy())
20   end
21 end
22
23 // Initialization and starting DFS from basic blocks with number 2
24 for (func_name, bb_num), block in blocks.items() do
25   if bb_num == 2 then
26     initial_isr_status ← track_isr_status(blocks).copy()
27     process_block(block, initial_isr_status)
28     for successor in block.successors do
29       dfs(successor, set(), initial_isr_status.copy(), [(func_name,
30         bb_num)])
31     end
32   end
33 end

```

---

Figure 3.5: Algorithm: DFS and Initialization

The **Depth-First Search** (DFS) function is recursively processing each block in a possible path of the RICFG. The set `visited_blocks` is used to avoid revisiting already visited blocks. If the block is already visited, the ISR status is updated with the stored ISR status for that block using the `merge_isr_statuses` function.

---

**Algorithm 8:** Merge ISR Statuses

---

**Data:** `isr_status1, isr_status2`

**Result:** Merged ISR status list

**Input:** Two lists of ISR statuses

**Output:** List of merged ISR statuses

```

1 merged_status ← empty list
2 for isr1, isr2 in zip(isr_status1, isr_status2) do
3   | merged_status.append(min(isr1, isr2))
4 end
5 return merged_status

```

---

Figure 3.6: Algorithm: Merge ISR Statuses

This function takes the worst case of the most enabled ISRs and uses this for further analysis of the path.

Unvisited blocks get added to the `visited_blocks` set and to the path list. After that, the ISR status gets updated to the current ISR status, and the function `process_block` is called to update the ISR status and track the shared resource accesses.

When the block is processed, the function checks for possible successors and recursively calls itself with the successor and the updated copy of `visited_blocks`, `current_isr_status`, and the path.

The first BB in a function is always the BB with number two in the generated cfg-files. To initialize the DFS, the BB number 2 is processed by the `process_block` function, and after that, the DFS function is called with the successor of the current block.

**Algorithm 9:** Check for Data Races**Data:** resource\_accesses**Result:** List of potential data races**Input:** Dictionary of resource accesses**Output:** List of potential data races

```

1 Function check_for_data_races():
2   for resource, accesses in resource_accesses.items() do
3     for i, (func1, bb_num1, access_type1, line_number1, isr_status1,
4       priority1) in enumerate(accesses) do
5       for j, (func2, bb_num2, access_type2, line_number2, isr_status2,
6         priority2) in enumerate(accesses) do
7         if i ≥ j then
8           Continue
9         end
10        if func1 ≠ func2 and (access_type1 == "write" or access_type2
11          == "write") and priority1 ≠ priority2 then
12          potential_data_races.append((resource, (func1, bb_num1,
13            access_type1, line_number1, isr_status1, priority1), (func2,
14              bb_num2, access_type2, line_number2, isr_status2,
15                priority2)))
16        end
17      end
18    end
19  end
20  // Call the function to check for data races
21 checkForDataRaces()

```

Figure 3.7: Algorithm: Check for Data Races

The function `check_for_data_races` identifies potential data races by comparing the pairs of data accesses that were initiated earlier. It iterates through all possible tuples of accesses. If a tuple is not within the same function, one of the two accesses is a write operation, and the priorities of both accesses are different, the pair is added to the list of possible data races. All the items in the list fulfill the conditions of a possible data race, which do not include the ISR status tracking. Since the ISR status tracking is the more complex part of the analysis, this makes sure to find all possible data races before filtering to minimize the number of false negatives.

## Filter Possible Data Races

---

**Algorithm 10:** Filter Data Races

---

**Data:** potential\_data\_races, isr\_enabling\_map  
**Result:** Filtered list of data races  
**Input:** List of potential data races, ISR enabling map  
**Output:** List of filtered data races

```

1 filtered_data_races ← empty list
2 seen_races ← empty set
3 for resource, access1, access2 in potential_data_races do
4     func1, bb_num1, access_type1, line_number1, isr_status1, priority1 ←
        access1
5     func2, bb_num2, access_type2, line_number2, isr_status2, priority2 ←
        access2
6     relevant_isr_disabled1 ← is_isr_disabled(isr_status1, func2) and not
        is_isr_enabled_by_another(isr_status1, func2)
7     relevant_isr_disabled2 ← is_isr_disabled(isr_status2, func1) and not
        is_isr_enabled_by_another(isr_status2, func1)
8     race_key ← frozenset(((func1, line_number1), (func2, line_number2)))
9     if not (relevant_isr_disabled1 or relevant_isr_disabled2) and race_key not
        in seen_races then
10         filtered_data_races.append((resource, access1, access2))
11         seen_races.add(race_key)
12     end
13 end
14 return filtered_data_races

```

---

Figure 3.8: Algorithm: Filter Data Races

The `filter_data_races` function takes the list of possible data races given by the `check_for_data_races` function and filters the racing pairs considering the ISR statuses of the involved ISRs. It takes the two accesses of a potential race and extracts the information that is saved in those accesses. After that, it uses two helper functions to determine the ISR statuses during the access.

---

**Algorithm 11:** Is ISR Disabled

---

**Data:** *isr\_status*, *func\_name***Result:** Boolean indicating if the ISR is disabled**Input:** List of ISR statuses, function name as a string**Output:** Boolean

```

1 isr_idx ← extract_isr_index(func_name)
2 if isr_idx is not None and isr_idx < length of isr_status then
3   | return isr_status[isr_idx] == 1
4 end
5 return False

```

---

Figure 3.9: Algorithm: Is ISR Disabled

The `is_isr_disabled` function checks if the bit corresponding to the ISR in the ISR status array is set to one. If so, the function returns true to the `filter_data_races` function, and if not, it returns false.

---

**Algorithm 12:** Is ISR Enabled by Another

---

**Data:** *isr\_status*, *func\_name*, *isr\_enabling\_map***Result:** Boolean indicating if the ISR is enabled by another function**Input:** List of ISR statuses, function name as a string, ISR enabling map**Output:** Boolean

```

1 isr_idx ← extract_isr_index(func_name)
2 if isr_idx is not None then
3   | for enabler_isr, enabled_isrs in isr_enabling_map.items() do
4     | enabler_idx ← extract_isr_index(enabler_isr)
5     | if enabler_idx is not None and not is_isr_disabled(isr_status,
6       | enabler_isr) then
7         | | if isr_idx in enabled_isrs then
8           | | | return True
9         | | end
10      | end
11 end
12 return False

```

---

Figure 3.10: Algorithm: Is ISR Enabled by Another

The `is_isr_enabled_by_another` function looks for possible activations of an ISR by another ISR. The `isr_enabling_map` was initiated and filled with information at the start of the `detect_data_races` function. This information is used in this function to determine if an ISR is enabled by another ISR that is enabled, to correctly handle racing pairs with these conditions.

## Parsing and Helper Functions

In this section the parsing of the input and the helper functions, which are called in the data race analysis, get explained.

---

### Algorithm 13: Parse Basic Blocks

---

**Data:** file\_path, shared\_resource\_names  
**Result:** blocks, function\_blocks  
**Input:** Path to file, List of shared resource names  
**Output:** Dictionary of BasicBlocks, Dictionary of function blocks

```

1 Initialize variables
2 lines ← read lines from file
3 for line in lines do
4   line ← trim(line)
5   if match function then
6     if bb_num and current_function then
7       | Add BasicBlock to blocks and function_blocks
8     end
9     current_function ← extract function name
10    bb_num ← None
11    Continue;
12  end
13  if match basic block then
14    if bb_num and current_function then
15      | Add BasicBlock to blocks and function_blocks
16    end
17    bb_num ← extract basic block number
18    Reset lists
19  end
20  for resource_name in shared_resource_names do
21    if resource_name in line then
22      if resource_name is written then
23        | Add write to shared_resources
24      end
25      else
26        | Add read to shared_resources
27      end
28    end
29  end
30  if line contains enable_isr or disable_isr then
31    | Add to enable_disable_calls
32  end
33  if match function call then
34    | Add function call to function_calls
35  end
36 end
37 if bb_num and current_function then
38 | Add BasicBlock to blocks and function_blocks
39 end
40 for line in lines do
41   line ← trim(line)
42   if match function then
43     current_function ← extract function name
44     bb_num ← None
45     Continue;
46   end
47   if line contains successors then
48     | Extract successors and update BasicBlock successors
49   end
50 end
51 return blocks, function_blocks

```

---

Figure 3.11: Algorithm: Parse Basic Blocks

The `parse_basic_block` function iterates two times through the code to extract all the important information of the code and save it in the BB items. In the first iteration of the code lines, the BBs get initiated with the BB number and the function it relates to. Additionally the information of shared resources, enable/disable calls of ISRs and function calls within the BB are added. In the second iteration the successors of the BB get added. A second iteration is used to ensure all the BBs items are initialized first to ensure a correct handling of the successors.

---

**Algorithm 14:** Determine Priority

---

**Data:** `function_name`

**Result:** Priority of the function

**Input:** Function name as a string

**Output:** Priority as an integer

```

1 match ← search for 'isr_?(+)' in function_name
2 if match is found then
3   | return integer value of matched group // Higher priority for lower ISR
   | number
4 end
5 else
6   | return infinity // Lower priority for non-ISR functions
7 end

```

---

Figure 3.12: Algorithm: Determine Priority

The `determine_priority` function is used to determine the priority of the function which is involved in a possible data race. Since one condition for a data race is two different priorities of the function this is an important check. The priority gets determined in the first place by differentiate between ISR functions and normal functions because ISRs always have a higher priority than non-ISR functions. Non-ISR functions have the priority infinity and ISRs get ordered by the number of it, while lower number ISRs have a higher priority than higher number ones.

---

**Algorithm 15:** Track ISR Status

---

**Data:** `blocks`

**Result:** List indicating ISR status initialized to zero

**Input:** Dictionary of BasicBlocks

**Output:** List of zeros representing the status of each ISR

```

1 isr_count ← count unique ISR function names in blocks
2 return list of zeros of length isr_count

```

---

Figure 3.13: Algorithm: Track ISR Status



---

**Algorithm 16:** Extract ISR Index

---

**Data:** `function_name`**Result:** Index of the ISR**Input:** Function name as a string**Output:** ISR index as an integer

```

1 match ← search for ISR pattern in function_name
2 if match then
3   | return integer value of matched group minus one
4 end
5 return None

```

---

Figure 3.14: Algorithm: Track ISR Status

---

**Algorithm 17:** Propagate Function Calls

---

**Data:** `blocks`, `function_blocks`**Result:** Updated blocks with propagated function call information**Input:** Dictionary of BasicBlocks, Dictionary of function blocks**Output:** Updated BasicBlocks with propagated resources and ISR calls

```

1 for each function in function_blocks do
2   | for each block in function's block list do
3     |   for each called_func, line_number in block's function calls do
4       |     if called_func in function_blocks then
5         |       | for each called_block in function_blocks[called_func] do
6           |         |   block.shared_resources.extend(called_block.shared_resources)
7           |         |   block.enable_disable_calls.extend(called_block.enable_disable_calls)
8         |       | end
9         |     end
10    |   end
11  | end
12 end

```

---

Figure 3.15: Algorithm: Propagate Function Calls

The `propagate_function_calls` function is handling the case of a function that calls another function. It checks for BBs items with a function call in it and adds the critical parts of the called function to the current BB to simulate a path through the called function and consider the shared resources and enable/disable calls of that function.



## 4 Evaluation

Done:

- Data Race Detection with all necessary criterias for a data race
- Informations of the cfg brought down to the minimum needed to analyze for data races
- Inter-Procedural checks of function calls
- Recursive traversal of the CFG using DFS
- Implementation of an ISR Status Array that updates thorough the traversal
- Considering ISR enabling ISRs

Not Done/ Future work:

- Pointer analysis
- Unnecessary nodes in CFG are empty but not deleted
- ISR enabling ISR with a depth of more than one are not considered (1 activates 2 activates 3 and 3 has data race)



## 5 Conclusion



# Bibliography

- [1] Wang, Y., Gao, F., Wang, L., Yu, T., Zhao, J., & Li, X. (2020). Automatic Detection, Validation, and Repair of Race Conditions in Interrupt-Driven Embedded Software. *IEEE Transactions on Software Engineering*.
- [2] Engler, D., & Ashcraft, K. (2003). RacerX: Effective, Static Detection of Race Conditions and Deadlocks. *ACM SIGOPS Operating Systems Review*.
- [3] Nikita Chopra, Rekha Pai, and Deepak D'Souza (2019). Data Races and Static Analysis for Interrupt-Driven Kernels
- [4] Flanagan, C., & Freund, S. N. (2009). FastTrack: Efficient and Precise Dynamic Race Detection. *ACM SIGPLAN Notices*.
- [5] R. Chen, Xiangying Guo, Y. Duan, B. Gu, Mengfei Yang (2011). Static Data Race Detection for Interrupt-Driven Embedded Software.
- [6] Muchnick, S. S. (1997). *Advanced Compiler Design and Implementation*. Morgan Kaufmann.
- [7] Adve, S. V., & Gharachorloo, K. (1996). Shared Memory Consistency Models: A Tutorial. *IEEE Computer*.
- [8] Herlihy, M., & Shavit, N. (2008). *The Art of Multiprocessor Programming*. Morgan Kaufmann.
- [9] K. Mehlhorn, P. Sanders (2008). *Algorithms and Data Structures The Basic Toolbox*.





# Attachments

```

1  class BasicBlock:
2  def __init__(function_name, number, shared_resources=[], successors
    =[], enable_disable_calls=[], code=[]):
3      self.function_name = function_name
4      self.number = number
5      self.shared_resources = shared_resources
6      self.successors = successors
7      self.enable_disable_calls = enable_disable_calls
8      self.code = code
9
10 def __repr__():
11 return ("BasicBlock(function_name={}, number={}, shared_resources={},
    "
12 "successors={}, enable_disable_calls={}, code={})".format(
13 self.function_name, self.number, self.shared_resources,
14 [succ.number for succ in self.successors], self.enable_disable_calls,
15 ' '.join(self.code)))
16
17 def parse_basic_blocks(file_path, shared_resource_names):
18     blocks = {}
19     current_function = None
20
21     with open(file_path, 'r') as file:
22         lines = file.readlines()
23
24     bb_num = None
25     shared_resources = []
26     enable_disable_calls = []
27     code_lines = []
28     line_number = 0
29
30     for line in lines:
31         line = line.strip()
32         line_number += 1
33
34         func_match = re.match(r';; Function (.+?) \(', line)
35         if func_match:
36             if bb_num is not None and current_function is not None:
37                 blocks[(current_function, bb_num)] = BasicBlock(
38                     current_function, bb_num, shared_resources, [], enable_disable_calls,
39                     code_lines)
40             current_function = func_match.group(1)
41             bb_num = None
42             continue
43
44         bb_match = re.match(r'<bb (\d+)>:', line)
45         if bb_match:
46             if bb_num is not None and current_function is not None:
47                 blocks[(current_function, bb_num)] = BasicBlock(
48                     current_function, bb_num, shared_resources, [], enable_disable_calls,
49                     code_lines)

```

```

48     bb_num = int(bb_match.group(1))
49     shared_resources = []
50     enable_disable_calls = []
51     code_lines = []
52
53     for resource_name in shared_resource_names:
54         if re.search(fr'\b{resource_name}\b', line):
55             if re.search(fr'\b{resource_name}\b\s*=', line):
56                 shared_resources.append((resource_name, 'write', line_number))
57             else:
58                 shared_resources.append((resource_name, 'read', line_number))
59
60         if 'enable_isr' in line or 'disable_isr' in line:
61             enable_disable_calls.append((line.strip(), line_number))
62
63     code_lines.append((line, line_number))
64
65     if bb_num is not None and current_function is not None:
66         blocks[(current_function, bb_num)] = BasicBlock(
67             current_function, bb_num, shared_resources, [], enable_disable_calls,
68             code_lines)
69
70     current_function = None
71     bb_num = None
72     for line in lines:
73         line = line.strip()
74
75         func_match = re.match(r';; Function (.+?) \(\'', line)
76         if func_match:
77             current_function = func_match.group(1)
78             bb_num = None
79             continue
80
81         if 'succs' in line:
82             succ_match = re.match(r';; (\d+) succs \{(.+?)\}', line)
83             if succ_match:
84                 bb_num = int(succ_match.group(1))
85                 succ_list = [int(succ.strip()) for succ in succ_match.group(2).split
86                             ()]
87                 if (current_function, bb_num) in blocks:
88                     blocks[(current_function, bb_num)].successors = [
89                         blocks[(current_function, succ)] for succ in succ_list if (
90                             current_function, succ) in blocks]
91
92     return blocks
93
94     def track_isr_status(blocks):
95         isr_count = len(set(block.function_name for block in blocks.values()
96                             if re.search(r'isr[_]?(\d+)', block.function_name)))
97         return [0] * isr_count
98
99     def extract_isr_index(function_name):
100         match = re.search(r'isr[_]?(\d+)', function_name)
101         if match:

```

```

98     return int(match.group(1)) - 1
99     return None
100
101     def detect_data_races(blocks):
102         potential_data_races = []
103         resource_accesses = defaultdict(list)
104         isr_enabling_map = defaultdict(set)
105
106         for block in blocks.values():
107             for call, line_number in block.enable_disable_calls:
108                 if 'enable_isr' in call:
109                     isr_idx_match = re.search(r'\((\d+)\)', call)
110                     if isr_idx_match:
111                         enabled_isr_idx = int(isr_idx_match.group(1)) - 1
112                         enabler_isr = block.function_name
113                         isr_enabling_map[enabler_isr].add(enabled_isr_idx)
114
115         def process_block(block, current_isr_status):
116             for line, line_number in block.code:
117                 if 'enable_isr' in line or 'disable_isr' in line:
118                     isr_idx_match = re.search(r'\((\d+)\)', line)
119                     if isr_idx_match:
120                         isr_idx = int(isr_idx_match.group(1)) - 1
121                         if "disable_isr" in line:
122                             if 0 <= isr_idx < len(current_isr_status):
123                                 current_isr_status[isr_idx] = 1
124                         elif "enable_isr" in line:
125                             if 0 <= isr_idx < len(current_isr_status):
126                                 current_isr_status[isr_idx] = 0
127
128
129             for resource_name, access_type, res_line_number in block.
130                 shared_resources:
131                 if res_line_number == line_number:
132                     resource_accesses[resource_name].append((block.function_name, block.
133                         number, access_type, line_number, current_isr_status.copy()))
134
135         def dfs(block, visited_blocks, current_isr_status, path):
136             if (block.function_name, block.number) in visited_blocks:
137                 return
138             visited_blocks.add((block.function_name, block.number))
139             path.append((block.function_name, block.number))
140
141             process_block(block, current_isr_status)
142
143             if not block.successors:
144                 pass
145             else:
146                 for successor in block.successors:
147                     dfs(successor, set(visited_blocks), current_isr_status.copy(), path.
148                         copy())
149
150         for (func_name, bb_num), block in blocks.items():

```

```

149     if bb_num == 2:
150         initial_isr_status = track_isr_status(blocks).copy()
151         process_block(block, initial_isr_status)
152         for successor in block.successors:
153             dfs(successor, set(), initial_isr_status.copy(), [(func_name, bb_num)
154                 ])
155
156     def check_for_data_races():
157         for resource, accesses in resource_accesses.items():
158             for i, (func1, bb_num1, access_type1, line_number1, isr_status1) in
159                 enumerate(accesses):
160                 for j, (func2, bb_num2, access_type2, line_number2, isr_status2) in
161                     enumerate(accesses):
162                     if i >= j:
163                         continue
164                     if func1 != func2 and (access_type1 == "write" or access_type2 == "
165                         write"):
166                         potential_data_races.append((resource, (func1, bb_num1, access_type1,
167                             line_number1, isr_status1),
168                             (func2, bb_num2, access_type2, line_number2, isr_status2)))
169
170     check_for_data_races()
171
172     def filter_data_races(potential_data_races):
173         filtered_data_races = []
174         for resource, access1, access2 in potential_data_races:
175             func1, bb_num1, access_type1, line_number1, isr_status1 = access1
176             func2, bb_num2, access_type2, line_number2, isr_status2 = access2
177
178     def is_isr_disabled(isr_status, func_name):
179         isr_idx = extract_isr_index(func_name)
180         if isr_idx is not None and isr_idx < len(isr_status):
181             return isr_status[isr_idx] == 1
182         return False
183
184     def is_isr_enabled_by_another(isr_status, func_name):
185         isr_idx = extract_isr_index(func_name)
186         if isr_idx is not None:
187             for enabler_isr, enabled_isrs in isr_enabling_map.items():
188                 enabler_idx = extract_isr_index(enabler_isr)
189                 if enabler_idx is not None and not is_isr_disabled(isr_status,
190                     enabler_isr):
191                     if isr_idx in enabled_isrs:
192                         return True
193             return False
194
195     relevant_isr_disabled1 = is_isr_disabled(isr_status1, func2) and not
196         is_isr_enabled_by_another(isr_status1, func2)
197     relevant_isr_disabled2 = is_isr_disabled(isr_status2, func1) and not
198         is_isr_enabled_by_another(isr_status2, func1)
199
200     if not (relevant_isr_disabled1 or relevant_isr_disabled2):

```

```
195     filtered_data_races.append((resource, access1, access2))
196
197     return filtered_data_races
198
199     filtered_data_races = filter_data_races(potential_data_races)
200
201     return filtered_data_races
202
203
204     shared_resource_input = input("Enter the names of shared resources,
205                                   separated by commas: ")
206     shared_resource_names = [name.strip() for name in
207                              shared_resource_input.split(',')]
208
209     file_path = input("Enter the file path: ").strip()
210     blocks = parse_basic_blocks(file_path, shared_resource_names)
211
212     data_races = detect_data_races(blocks)
213
214     print("Detected Data Races:")
215     for resource, access1, access2 in data_races:
216         print(f"Resource: {resource}")
217         print(f"  Access 1: Function {access1[0]} (BB {access1[1]}), {access1[2]}, Line {access1[3]}, ISR Status: {access1[4]}")
218         print(f"  Access 2: Function {access2[0]} (BB {access2[1]}), {access2[2]}, Line {access2[3]}, ISR Status: {access2[4]}")
219         print()
```